

An Experimental Comparison of Concurrent Data Structures

Student: Mark Gibson Supervisor: Dr. David Gregg

Background

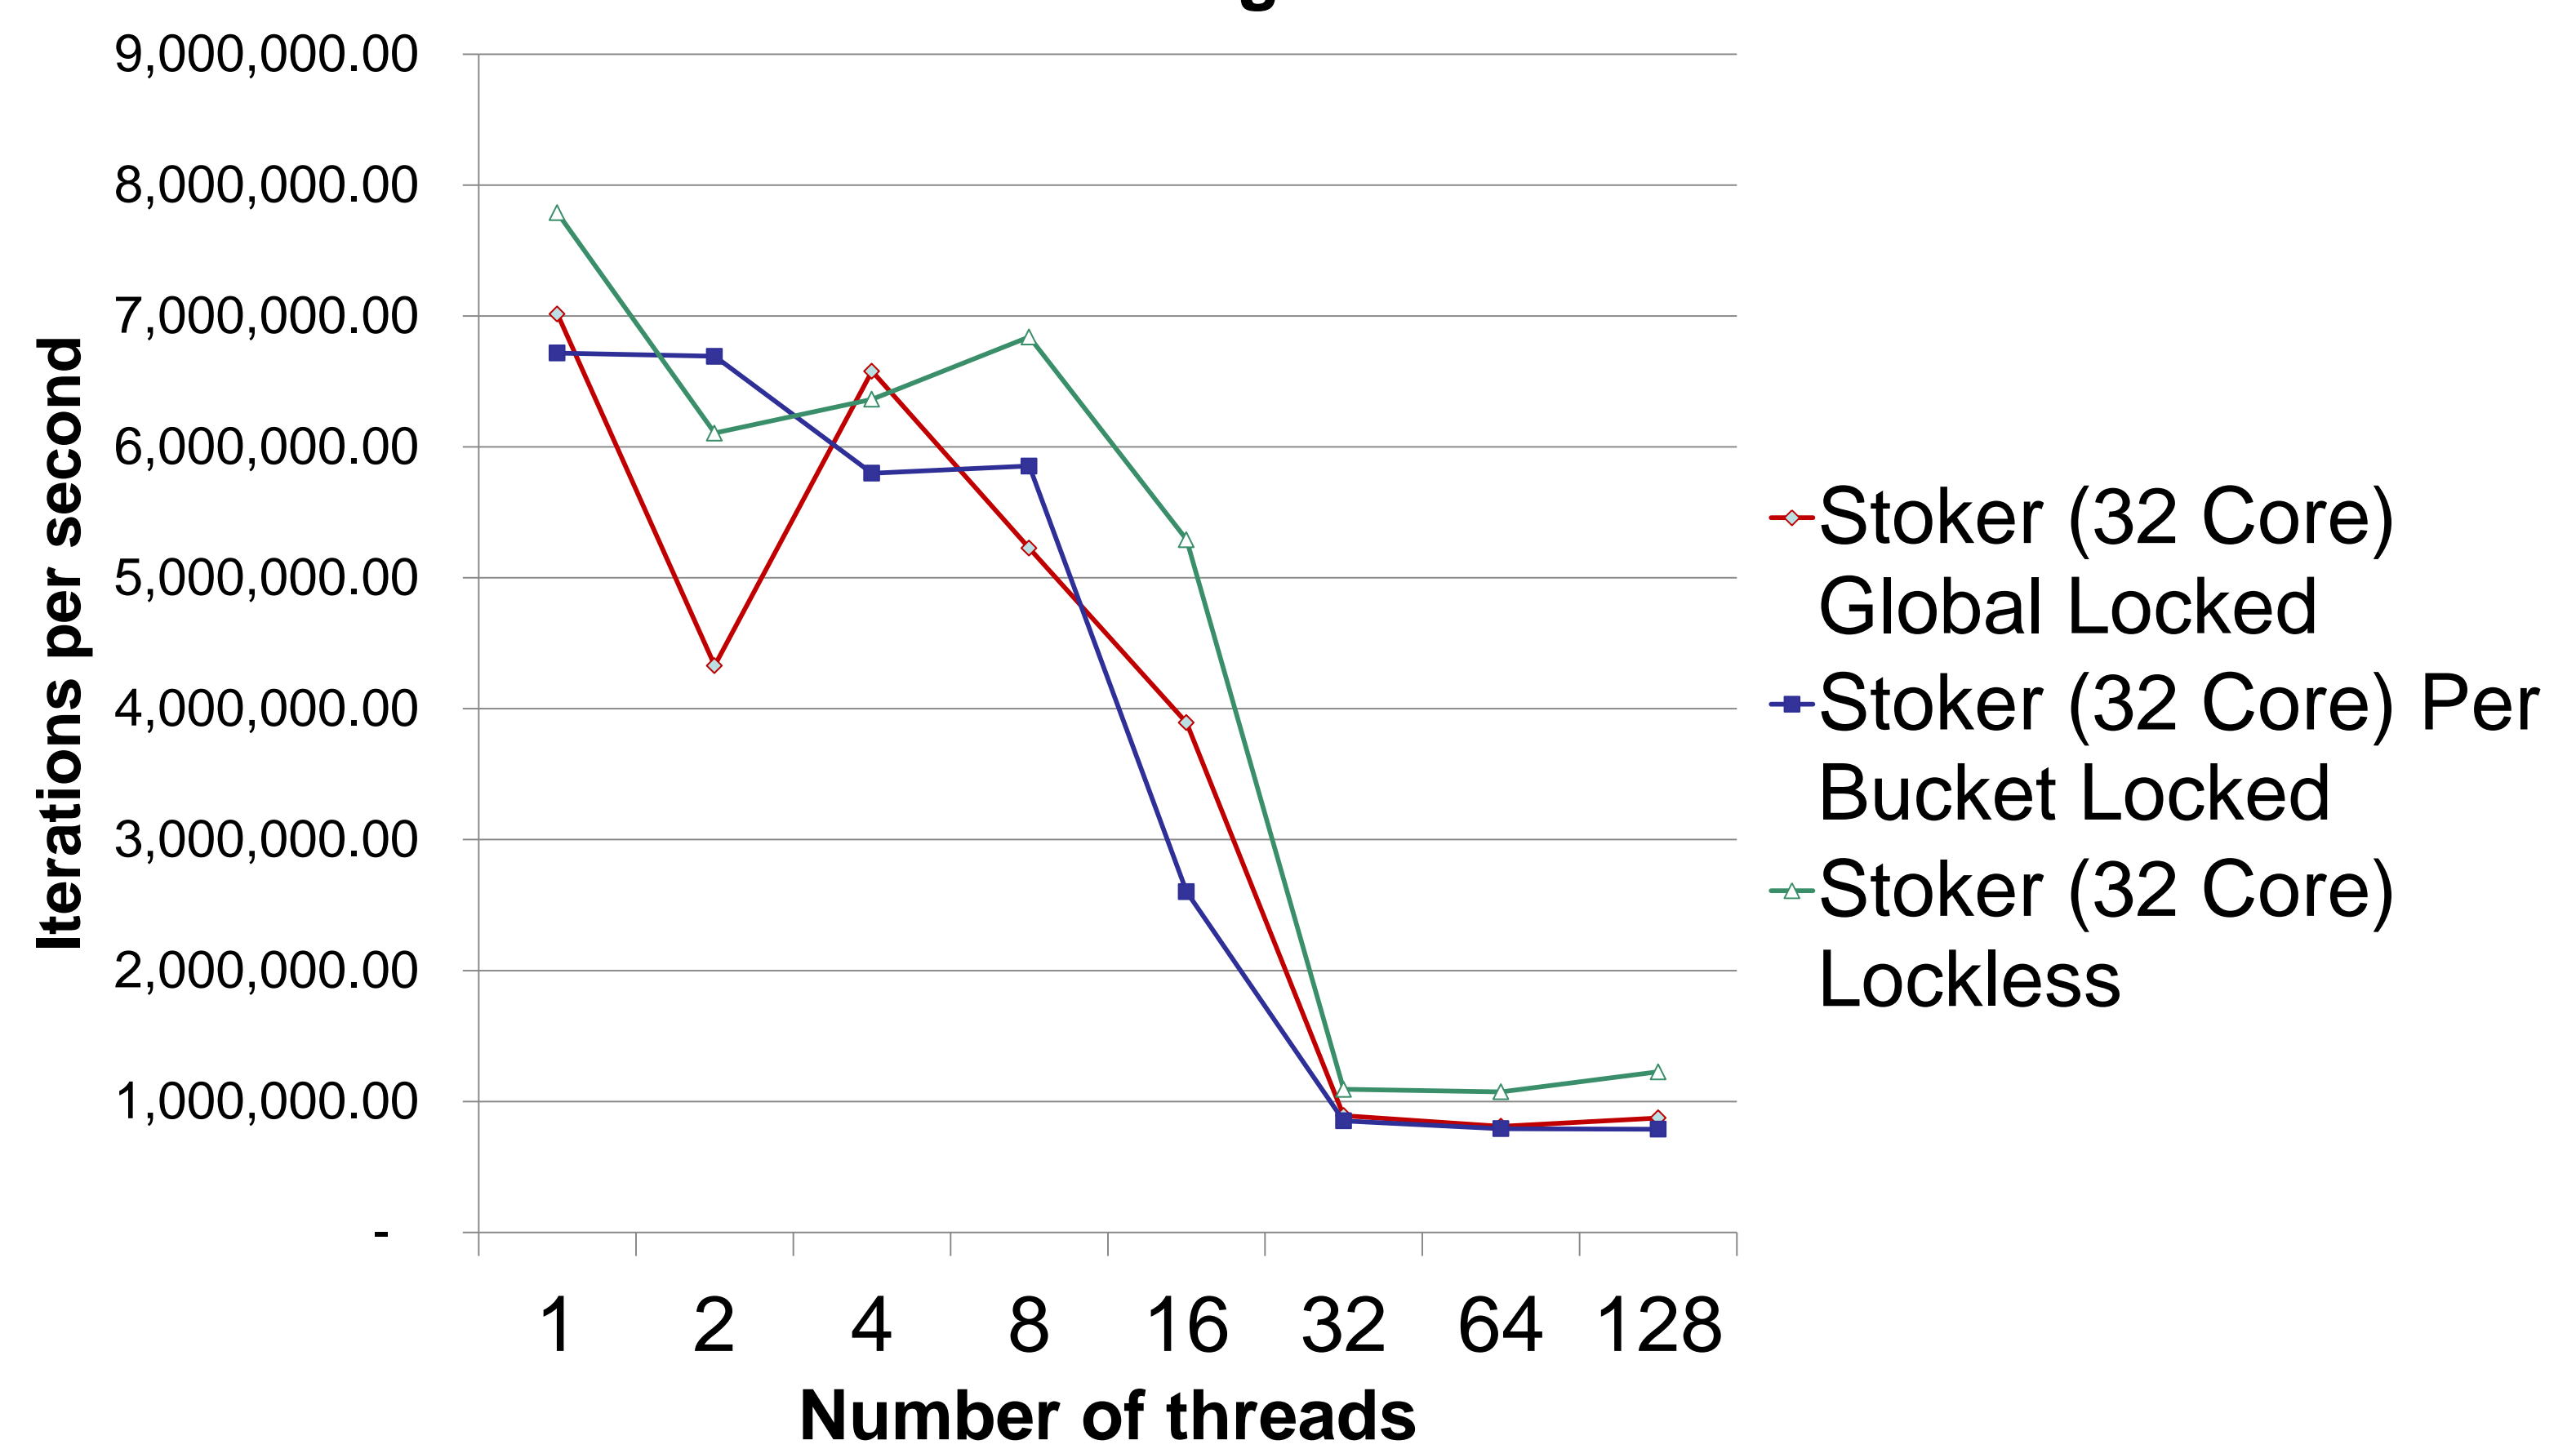
- Concurrent Data Structures
 - Designed for access by multiple threads
 - Potential for high scalability
- Atomic Instructions
 - Either complete fully or not at all
 - Used to implement locked & lockless algorithms
- Locks
 - Use mutexes to acquire/release a lock
 - Blocks threads that don't have the lock
- Lockless
 - Use atomic instructions
 - Guaranteed system throughput

Approach

- Implemented three concurrent data structures
 - Ring Buffer
 - Linked List
 - Hash Table
- Lock Variations
 - Pthread Mutex
 - Test-and-Set (3 Variations)
 - Ticket (2 Variations)
- Lockless Variations
 - Compare-and-Swap
 - Fetch-and-Add
- 13 Variations for each Data Structure
 - 12 Locks, 1 Lockless
 - Compared each on 3 systems


Evaluation

Hash Table; Stoker; Pthread Mutex lock vs lockless; Key Range 128



	Global Lock	Bucket Lock	Lockless
CPU Cycles	1.13x10 ¹²	9.86x10 ¹¹	1.09x10 ¹²
Stalled Front End	1.06x10 ¹²	9.16x10 ¹¹	1.02x10 ¹²
Stalled Back End	6.89x10 ¹¹	5.43x10 ¹¹	6.27x10 ¹¹

Singly Linked Buffer; Cube; Mutex Lock vs Lockless; Key Range 128



	Cube Lockless	Cube Locked
CPU Cycles	1.66x10 ¹⁰	3.70x10 ¹⁰
Stalled Front End	8.06x10 ⁹	2.81x10 ¹⁰
Stalled Back End	4.29x10 ⁹	1.92x10 ¹⁰

Conclusion

- The project was successful in performing an in-depth comparison of 3 data structures
- Lockless algorithms were mostly faster
 - Some exceptions
- The linked list had the best performance gain
 - Using Lockless algorithm versus Lock
- The hash table had the worst performance
 - Outperformed by the locked variations